

# demo8\_workbook\_MASTER

October 15, 2022

## 1 MASTER wk8 Demo - Advanced Spark - DataFrames and Spark SQL

MIDS w261: Machine Learning at Scale | UC Berkeley School of Information | Fall 2022

So far we've been using Spark's low level APIs. In particular, we've been using the RDD (Resilient Distributed Datasets) API to implement Machine Learning algorithms from scratch. This week we're going to take a look at how Spark is used in a production setting. We'll look at DataFrames, SQL, and UDFs (User Defined Functions). As discussed previously, we still need to understand the internals of Spark and MapReduce in general to write efficient and scalable code.

In class today we'll get some practice working with larger data sets in Spark. We'll start with an introduction to efficiently storing data and approach a large dataset for analysis. After that we'll discuss a ranking problem which was covered in Chapter 6 of the High Performance Spark book and how we can apply that to our problem. We'll follow up with a discussion on things that could be done to make this more efficient. \* ... **describe** differences between data serialization formats. \* ... **choose** a data serialization format based on use case. \* ... **describe** DataFrames API, GroupBy and *Spark SQL*. \* ... **describe** and **create** a data pipeline for analysis. \* ... **use** a user defined function (UDF). \* ... **understand** feature engineering and aggregations in Spark.

**Additional Resources:** Writing performant code in Spark requires a lot of thought. Holden's High Performance Spark book covers this topic very well. In addition, Spark - The Definitive Guide, by Bill Chambers and Matei Zaharia, provides some recent developments.

```
[1]: ## Imports
import re
import json
import numpy as np
import pandas as pd
import warnings
warnings.filterwarnings('ignore')

from pyspark.sql import SparkSession

app_name = "week8_demo"
master = "local[*]"
spark = SparkSession\
    .builder\
```

```

    .appName(app_name)\
    .master(master)\
    .config("spark.ui.port","42229")\
    .getOrCreate()
sc = spark.sparkContext

## Change the working directory
!cd /media/notebooks/student-workspace/LiveSessionMaterials/wk08Demo_DataFrames

```

```

:: loading settings :: url = jar:file:/usr/lib/spark/jars/ivy-2.4.0.jar!/org/apache/ivy/core/settings/ivysettings.xml

```

Ivy Default Cache set to: /root/.ivy2/cache

The jars for the packages stored in: /root/.ivy2/jars

graphframes#graphframes added as a dependency

org.apache.spark#spark-avro\_2.12 added as a dependency

```

:: resolving dependencies :: org.apache.spark#spark-submit-parent-2825d7bc-faf1-4b84-8258-67801c3c6190;1.0

```

```

  confs: [default]

```

```

  found graphframes#graphframes;0.8.2-spark3.1-s_2.12 in spark-packages

```

```

  found org.slf4j#slf4j-api;1.7.16 in central

```

```

  found org.apache.spark#spark-avro_2.12;3.1.3 in central

```

```

  found org.spark-project.spark#unused;1.0.0 in central

```

```

:: resolution report :: resolve 325ms :: artifacts dl 7ms

```

```

  :: modules in use:

```

```

  graphframes#graphframes;0.8.2-spark3.1-s_2.12 from spark-packages in

```

```

[default]

```

```

  org.apache.spark#spark-avro_2.12;3.1.3 from central in [default]

```

```

  org.slf4j#slf4j-api;1.7.16 from central in [default]

```

```

  org.spark-project.spark#unused;1.0.0 from central in [default]

```

```

-----
|               |             modules             ||   artifacts   | | | | |
|               | number| search|dwnlded|evicted|| number|dwnlded|
|               |-----|-----|-----|-----||-----|-----|
|               | 4    | 0    | 0    | 0    || 4    | 0    |
|               |-----|-----|-----|-----||-----|-----|

```

```

:: retrieving :: org.apache.spark#spark-submit-parent-2825d7bc-faf1-4b84-8258-67801c3c6190

```

```

  confs: [default]

```

```

  0 artifacts copied, 4 already retrieved (0kB/8ms)

```

Setting default log level to "WARN".

To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).

22/10/08 15:30:25 INFO org.apache.spark.SparkEnv: Registering MapOutputTracker

22/10/08 15:30:25 INFO org.apache.spark.SparkEnv: Registering BlockManagerMaster

22/10/08 15:30:25 INFO org.apache.spark.SparkEnv: Registering

BlockManagerMasterHeartbeat

22/10/08 15:30:25 INFO org.apache.spark.SparkEnv: Registering

OutputCommitCoordinator

/bin/sh: 1: cd: can't cd to /media/notebooks/student-workspace/LiveSessionMaterials/wk08Demo\_DataFrames

```
[3]: ## Load the data
data = spark.read.format('bigquery') \
    .option('table', 'bigquery-public-data:samples.gsod') \
    .load()
```

## 1.1 DataFrames API

Let's showcase some of the important methods that we have available when working with DataFrames

```
[3]: ## show
data.show()
```

22/06/20 15:32:50 WARN org.apache.spark.sql.catalyst.util.package: Truncated the string representation of a plan since it was too large. This behavior can be adjusted by setting 'spark.sql.debug.maxToStringFields'.

[Stage 0:> (0 + 1) / 1]

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|station_number|wban_number|year|month|day|
mean_temp|num_mean_temp_samples|mean_dew_point|num_mean_dew_point_samples|me
an_sealevel_pressure|num_mean_sealevel_pressure_samples|mean_station_pressure|nu
m_mean_station_pressure_samples|mean_visibility|num_mean_visibility_samples|
mean_wind_speed|num_mean_wind_speed_samples|max_sustained_wind_speed|max_gust_wi
nd_speed|max_temperature|max_temperature_explicit|min_temperature|min_tempera
ture_explicit|total_precipitation|snow_depth|fog|rain|snow|
hail|thunder|tornado|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|39730|99999|1929|10|20|52.79999923706055|
4|45.5|4|null|
null|null|null|6.199999809265137|
4|21.200000762939453|4|29.899999618530273|
```

null	50.0	false	null
null	0.0	null false false false false	false  false
	33110	99999 1929  12  18	47.5
4	44.0	4	null
null	null	null	2.5
4	11.0	4	13.0
null	45.0	false	null
null	null	null false false false false	false  false
	37770	99999 1931  4  24	50.20000076293945
4	44.29999923706055	4	null
null	null	null	5.900000095367432
4	12.300000190734863	4	18.100000381469727
null	45.0	false	null
null	null	null false false false false	false  false
	726810	24131 1931  6  23	65.0999984741211
24	41.5	8	null
null	null	null	48.29999923706055
24	7.199999809265137	24	11.100000381469727
null	53.400001525878906	true	null
null	0.0	null false false false false	false  false
	726810	24131 1931  3  2	42.79999923706055
24	31.5	8	null
null	null	null	72.9000015258789
24	2.299999952316284	24	4.099999904632568
null	32.400001525878906	true	null
null	0.0	null false false false false	false  false
	726810	24131 1931  9  17	67.0
24	40.5	8	null
null	null	null	33.79999923706055
24	2.4000000953674316	24	6.0
null	51.29999923706055	true	null
null	0.0	null false false false false	false  false
	726810	24131 1931  8  7	68.4000015258789
24	37.20000076293945	8	null
null	null	null	27.899999618530273
24	3.5	24	7.0
null	52.29999923706055	true	null
null	0.0	null false false false false	false  false
	726810	24131 1932  7  14	64.0999984741211
24	54.099998474121094	8	null
null	null	null	41.0
24	4.199999809265137	24	8.899999618530273
null	55.400001525878906	true	null
null	null	null false false false false	false  false
	726810	24131 1932  10  23	41.099998474121094
24	31.0	8	null
null	null	null	41.0
24	4.300000190734863	24	15.899999618530273

null	35.400001525878906					true		null
null		726810	24131	1932	1	5	24.600000381469727	4.099999904632568
24	21.100000381469727					8		null
null								14.800000190734863
24	3.200000047683716					24		4.099999904632568
null	21.399999618530273					true		null
null		726815	24106	1932	8	27	71.0	
24								15.0
null								28.600000381469727
24		8.0				24		
null	62.400001525878906					true		null
null		0.0				null	false	false
24	41.70000076293945					8		32.5
null								9.899999618530273
24	4.099999904632568					24		9.899999618530273
null	53.400001525878906					true		null
null		0.0				null	false	false
24	46.599998474121094					8		22.600000381469727
null								8.0
24	5.300000190734863					24		
null	46.400001525878906					true		null
null						null	false	false
4		370310	99999	1933	10	17	55.29999923706055	
null								3.0
4	16.799999237060547					4		18.100000381469727
null		45.0				false		null
null		0.0				null	false	false
4		292310	99999	1933	12	1	-11.0	
null								4.0
4		5.0				4		8.899999618530273
null		-27.0				true		null
null						null	false	false
4		370310	99999	1933	6	17	62.70000076293945	
null								1.899999976158142
null		1.0				4		
null		55.0				true		null
null		0.0				null	false	false
4		239330	99999	1933	6	4	70.5	
null								1.899999976158142
4		5.0				4		8.899999618530273

```

null|          59.0|          false|          null|
null|          0.0|          null| true| true| true| true|  true|  true|
|      282750|      99999|1933|      1|  7|          -5.0|
4|          null|          null|          null|          null|
null|          null|          null|          null|3.4000000953674316|
4| 3.200000047683716|          4|          8.899999618530273|
null|          -22.0|          true|          null|
null|          0.0|          null|false|false|false|false| false| false|
|      292310|      99999|1933|      3| 17|-10.300000190734863|
4|          null|          null|          null|          null|
null|          null|          null|          null| 1.100000023841858|
4|          13.5|          4|          18.100000381469727|
null|          -36.0|          true|          null|
null|          0.0|          null|false|false|false|false| false| false|
|      370310|      99999|1933|      4| 23|          65.0|
4|          null|          null|          null|          null|
null|          null|          null|          null|          null|
null|          4.5|          4|          8.899999618530273|
null|          52.0|          false|          null|
null|          0.0|          null|false|false|false|false| false| false|
+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+
only showing top 20 rows

```

Here we see `.show()`, a method that works similarly to Pandas `.head()`. You can observe that the DataFrame is stored in text, that way it's easier to distribute throughout the different executors. If you want to better display the results, we can transform the output using `.limit(n)` to a Pandas DataFrame

```
[4]: data.limit(10).toPandas().head()
```

```

[4]:  station_number  wban_number  year  month  day  mean_temp  \
0           39730         99999  1929    10    20  52.799999
1           33110         99999  1929    12    18  47.500000
2           37770         99999  1931     4    24  50.200001
3           726810        24131  1931     6    23  65.099998
4           726810        24131  1931     3     2  42.799999

   num_mean_temp_samples  mean_dew_point  num_mean_dew_point_samples  \
0                      4          45.500000                      4
1                      4          44.000000                      4

```

2	4	44.299999	4
3	24	41.500000	8
4	24	31.500000	8

	mean_sealevel_pressure	...	min_temperature	min_temperature_explicit	\
0	NaN	...	NaN	None	
1	NaN	...	NaN	None	
2	NaN	...	NaN	None	
3	NaN	...	NaN	None	
4	NaN	...	NaN	None	

	total_precipitation	snow_depth	fog	rain	snow	hail	thunder	\
0	0.0	NaN	False	False	False	False	False	
1	NaN	NaN	False	False	False	False	False	
2	NaN	NaN	False	False	False	False	False	
3	0.0	NaN	False	False	False	False	False	
4	0.0	NaN	False	False	False	False	False	

	tornado
0	False
1	False
2	False
3	False
4	False

[5 rows x 31 columns]

This is a public dataset from NOAA, regarding weather stations across the United States. It has a total of 31 columns.

Another important command is `.printSchema()` to check columns names and what type of data is stored on it

```
[5]: data.printSchema()
```

```
root
|-- station_number: long (nullable = false)
|-- wban_number: long (nullable = true)
|-- year: long (nullable = false)
|-- month: long (nullable = false)
|-- day: long (nullable = false)
|-- mean_temp: double (nullable = true)
|-- num_mean_temp_samples: long (nullable = true)
|-- mean_dew_point: double (nullable = true)
|-- num_mean_dew_point_samples: long (nullable = true)
|-- mean_sealevel_pressure: double (nullable = true)
|-- num_mean_sealevel_pressure_samples: long (nullable = true)
|-- mean_station_pressure: double (nullable = true)
```

```

|-- num_mean_station_pressure_samples: long (nullable = true)
|-- mean_visibility: double (nullable = true)
|-- num_mean_visibility_samples: long (nullable = true)
|-- mean_wind_speed: double (nullable = true)
|-- num_mean_wind_speed_samples: long (nullable = true)
|-- max_sustained_wind_speed: double (nullable = true)
|-- max_gust_wind_speed: double (nullable = true)
|-- max_temperature: double (nullable = true)
|-- max_temperature_explicit: boolean (nullable = true)
|-- min_temperature: double (nullable = true)
|-- min_temperature_explicit: boolean (nullable = true)
|-- total_precipitation: double (nullable = true)
|-- snow_depth: double (nullable = true)
|-- fog: boolean (nullable = true)
|-- rain: boolean (nullable = true)
|-- snow: boolean (nullable = true)
|-- hail: boolean (nullable = true)
|-- thunder: boolean (nullable = true)
|-- tornado: boolean (nullable = true)

```

```

[7]: %%time
    ## To look how many data points, we can use the command .count()
    print(f"Number of rows is {data.count()} and number of columns is {len(data.
    ↪columns)}")

```

[Stage 5:> (0 + 4) / 4]

```

Number of rows is 114420316 and number of columns is 31
CPU times: user 2.71 ms, sys: 2.83 ms, total: 5.54 ms
Wall time: 1.53 s

```

114 million rows! Try to fit that into a Pandas DataFrame!. Now let's check how can we filter our dataframe and how can we create new columns.

We need to lever a very important set of Spark built-in functions from `pyspark.sql.functions`, typically called F functions

```

[27]: # Using built-in Spark functions are always more efficient
from pyspark.sql import types
import pyspark.sql.functions as F

## Let's create a new column called time
data = data.withColumn("time",
                        F.concat(F.col("year"),
                                F.lit("-"), F.col("month"),
                                F.lit("-"), F.col("day")) \
                        .cast(types.TimestampType()))

```



```
data.printSchema()
```

```
root
|-- station_number: long (nullable = false)
|-- wban_number: long (nullable = true)
|-- year: long (nullable = false)
|-- month: long (nullable = false)
|-- day: long (nullable = false)
|-- mean_temp: double (nullable = true)
|-- num_mean_temp_samples: long (nullable = true)
|-- mean_dew_point: double (nullable = true)
|-- num_mean_dew_point_samples: long (nullable = true)
|-- mean_sealevel_pressure: double (nullable = true)
|-- num_mean_sealevel_pressure_samples: long (nullable = true)
|-- mean_station_pressure: double (nullable = true)
|-- num_mean_station_pressure_samples: long (nullable = true)
|-- mean_visibility: double (nullable = true)
|-- num_mean_visibility_samples: long (nullable = true)
|-- mean_wind_speed: double (nullable = true)
|-- num_mean_wind_speed_samples: long (nullable = true)
|-- max_sustained_wind_speed: double (nullable = true)
|-- max_gust_wind_speed: double (nullable = true)
|-- max_temperature: double (nullable = true)
|-- max_temperature_explicit: boolean (nullable = true)
|-- min_temperature: double (nullable = true)
|-- min_temperature_explicit: boolean (nullable = true)
|-- total_precipitation: double (nullable = true)
|-- snow_depth: double (nullable = true)
|-- fog: boolean (nullable = true)
|-- rain: boolean (nullable = true)
|-- snow: boolean (nullable = true)
|-- hail: boolean (nullable = true)
|-- thunder: boolean (nullable = true)
|-- tornado: boolean (nullable = true)
|-- time: timestamp (nullable = true)
```

```
[9]: ## If you want to select one or a set of columns, we can use the select method
data.select('time').show(5)
```

```
+-----+
|           time|
+-----+
|1929-10-20 00:00:00|
|1929-12-18 00:00:00|
|1931-04-24 00:00:00|
|1931-06-23 00:00:00|
```

```
|1931-03-02 00:00:00|
+-----+
only showing top 5 rows
```

```
[10]: data.select(['time', 'tornado']).show(5)
```

```
+-----+-----+
|           time|tornado|
+-----+-----+
|1929-10-20 00:00:00|  false|
|1929-12-18 00:00:00|  false|
|1931-04-24 00:00:00|  false|
|1931-06-23 00:00:00|  false|
|1931-03-02 00:00:00|  false|
+-----+-----+
only showing top 5 rows
```

```
[11]: ## If you want any row, we can take
data.take(1)
```

```
[11]: [Row(station_number=39730, wban_number=99999, year=1929, month=10, day=20,
mean_temp=52.79999923706055, num_mean_temp_samples=4, mean_dew_point=45.5,
num_mean_dew_point_samples=4, mean_sealevel_pressure=None,
num_mean_sealevel_pressure_samples=None, mean_station_pressure=None,
num_mean_station_pressure_samples=None, mean_visibility=6.199999809265137,
num_mean_visibility_samples=4, mean_wind_speed=21.200000762939453,
num_mean_wind_speed_samples=4, max_sustained_wind_speed=29.899999618530273,
max_gust_wind_speed=None, max_temperature=50.0, max_temperature_explicit=False,
min_temperature=None, min_temperature_explicit=None, total_precipitation=0.0,
snow_depth=None, fog=False, rain=False, snow=False, hail=False, thunder=False,
tornado=False, time=datetime.datetime(1929, 10, 20, 0, 0))]
```

Each Row of the DataFrame is a Row which is similar to a dictionary, you can reference each element of the Row using the key. Now, also notice that the output of `take` is a list, so you need to index the list first

```
[12]: ## let's get the station_number only
data.take(1)[0]['station_number']
```

```
[12]: 39730
```

```
[13]: ## Let's check now how to filter data using another weather station data
stations = spark.read.format('bigquery') \
    .option('table', 'bigquery-public-data:noaa_gsod.stations') \
    .load()
```

```
[14]: ## Let's filter only US based stations
# Let's filter for just the US since this is a US based dataset
stations_us = stations.filter(F.col('Country')== 'US')

print(f'Total stations are {stations.count()}, total US stations are_
↳{stations_us.count()}')
```

Total stations are 29590, total US stations are 7161

```
[15]: %%time
## Finally, we can describe our dataset using the describe command, similar to_
↳Pandas
## Let's select just a few columns
keep_columns = ['station_number', 'mean_temp', 'thunder',_
↳'mean_sealevel_pressure']
data.select(keep_columns).describe().show()
```

[Stage 18:=====>

(2 + 2) / 4]

summary	station_number	mean_temp	mean_sealevel_pressure
count	114420316	114420316	86731897
mean	507199.9578684261	52.122209996445854	1014.8442087525018
stddev	298384.12645319354	24.222342560059765	9.38153057246377
min	8209	-118.5	900.0
max	999999	110.0	1079.699951171875

CPU times: user 11.1 ms, sys: 1.66 ms, total: 12.8 ms

Wall time: 18.9 s

## 2 Data Types

I highly recommend reading this article [Format Wars](#) which covered the characteristics, structure, and differences between raw text, sequence, Avro, Parquet, and ORC data serializations.

There were several points discussed:

- Human Readable
- Row vs Column Oriented
- Read vs Write performance
- Appendable
- Splittable
- Metadata storage

We have 4 data types below

- Compressed CSV
- Parquet
- Avro
- CSV

Of these 3 are row oriented and 1 is column oriented. We have over 100M rows and 31 columns. Columnar compression should do fairly well in this scenerio.

```
[4]: # Access staging bucket and see whats there
import os
GCS_LOCATION = os.getenv('DATA_BUCKET')
GCS_LOCATION
```

STAGING\_BUCKET location: gs://dataproc-staging-us-central1-1077780374322-lwjldfuu/

```
[ ]: %%time
!gsutil rm -r {GCS_LOCATION}datagzip
data.write.option("compression","gzip").csv(f'{GCS_LOCATION}datagzip')
!gsutil du -sh {GCS_LOCATION}datagzip/*
```

Removing gs://dataproc-staging-us-central1-1077780374322-lwjldfuu/datagzip/#1655740918347358...  
/ [1 objects]  
Operation completed over 1 objects.

3.08 GiB gs://dataproc-staging-us-central1-1077780374322-lwjldfuu/datagzip/\*  
CPU times: user 179 ms, sys: 62.5 ms, total: 242 ms  
Wall time: 9min 53s

```
[ ]: %%time
!gsutil rm -r {GCS_LOCATION}dataparquet
data.write.format("parquet").save(f'{GCS_LOCATION}dataparquet')
!gsutil du -sh {GCS_LOCATION}dataparquet/*
```

Removing gs://dataproc-staging-us-central1-1077780374322-lwjldfuu/dataparquet/#1655741566985968...  
Removing gs://dataproc-staging-us-central1-1077780374322-lwjldfuu/dataparquet/\_SUCCESS#1655741567151993...  
Removing gs://dataproc-staging-us-central1-1077780374322-lwjldfuu/dataparquet/part-00000-ecbf1a52-7da7-4857-9fdd-f48c9c03610f-c000.snappy.parquet#1655741565407991...  
Removing gs://dataproc-staging-us-central1-1077780374322-lwjldfuu/dataparquet/part-00001-ecbf1a52-7da7-4857-9fdd-f48c9c03610f-c000.snappy.parquet#1655741565274501...  
/ [4 objects]  
==> NOTE: You are performing a sequence of gsutil operations that may run significantly faster if you instead use gsutil -m rm ... Please

see the `-m` section under "gsutil help options" for further information about when `gsutil -m` can be advantageous.

Removing gs://dataproc-staging-us-central1-1077780374322-lwjldfuu/dataparquet/part-00002-ecbf1a52-7da7-4857-9fdd-f48c9c03610f-c000.snappy.parquet#1655741566582027...

Removing gs://dataproc-staging-us-central1-1077780374322-lwjldfuu/dataparquet/part-00003-ecbf1a52-7da7-4857-9fdd-f48c9c03610f-c000.snappy.parquet#1655741561627949...

/ [6 objects]

Operation completed over 6 objects.

1.72 GiB gs://dataproc-staging-us-central1-1077780374322-lwjldfuu/dataparquet/\*  
CPU times: user 142 ms, sys: 40.8 ms, total: 183 ms  
Wall time: 3min 54s

```
[ ]: %%time
!gsutil rm -r {GCS_LOCATION}dataavro
data.write.format("avro").save(f'{GCS_LOCATION}dataavro')
!gsutil du -sh {GCS_LOCATION}dataavro/*
```

Removing gs://dataproc-staging-us-central1-1077780374322-lwjldfuu/dataavro/#1655742062236332...

Removing gs://dataproc-staging-us-central1-1077780374322-lwjldfuu/dataavro/\_SUCCESS#1655742970570332...

Removing gs://dataproc-staging-us-central1-1077780374322-lwjldfuu/dataavro/part-00000-2441faa1-1680-4894-920a-660b398f8048-c000.avro#1655742048661975...

Removing gs://dataproc-staging-us-central1-1077780374322-lwjldfuu/dataavro/part-00000-76afd8a5-f31e-4b87-a48d-f6c7f22701ea-c000.avro#1655742548203079...

/ [4 objects]

==> NOTE: You are performing a sequence of `gsutil` operations that may run significantly faster if you instead use `gsutil -m rm ...`. Please see the `-m` section under "gsutil help options" for further information about when `gsutil -m` can be advantageous.

Removing gs://dataproc-staging-us-central1-1077780374322-lwjldfuu/dataavro/part-00001-2441faa1-1680-4894-920a-660b398f8048-c000.avro#1655742060703836...

Removing gs://dataproc-staging-us-central1-1077780374322-lwjldfuu/dataavro/part-00001-76afd8a5-f31e-4b87-a48d-f6c7f22701ea-c000.avro#1655742549208485...

Removing gs://dataproc-staging-us-central1-1077780374322-lwjldfuu/dataavro/part-00002-2441faa1-1680-4894-920a-660b398f8048-c000.avro#1655742061834562...

Removing gs://dataproc-staging-us-central1-1077780374322-lwjldfuu/dataavro/part-00002-76afd8a5-f31e-4b87-a48d-f6c7f22701ea-c000.avro#1655742970078689...

Removing gs://dataproc-staging-us-central1-1077780374322-lwjldfuu/dataavro/part-00003-2441faa1-1680-4894-920a-660b398f8048-c000.avro#1655742059267771...

Removing gs://dataproc-staging-us-central1-1077780374322-lwjldfuu/dataavro/part-

```
00003-76afd8a5-f31e-4b87-a48d-f6c7f22701ea-c000.avro#1655742959591683...
/ [10 objects]
Operation completed over 10 objects.
```

```
4.66 GiB      gs://dataproc-staging-us-central1-1077780374322-lwjldfuu/dataavro/*
CPU times: user 179 ms, sys: 49 ms, total: 228 ms
Wall time: 7min 55s
```

```
[ ]: %%time
!gsutil rm -r {GCS_LOCATION}datacsv
data.write.csv(f'{GCS_LOCATION}datacsv')
!gsutil du -sh {GCS_LOCATION}datacsv/*
```

```
Removing gs://dataproc-staging-us-
central1-1077780374322-lwjldfuu/datacsv/#1655742810311463...
Removing gs://dataproc-staging-us-
central1-1077780374322-lwjldfuu/datacsv/_SUCCESS#1655742810479233...
Removing gs://dataproc-staging-us-central1-1077780374322-lwjldfuu/datacsv/part-0
0000-184cd7ba-1147-494d-ae8-e8da0bb34a25-c000.csv#1655742431834302...
Removing gs://dataproc-staging-us-central1-1077780374322-lwjldfuu/datacsv/part-0
0001-184cd7ba-1147-494d-ae8-e8da0bb34a25-c000.csv#1655742442219329...
/ [4 objects]
==> NOTE: You are performing a sequence of gsutil operations that may
run significantly faster if you instead use gsutil -m rm ... Please
see the -m section under "gsutil help options" for further information
about when gsutil -m can be advantageous.
```

```
Removing gs://dataproc-staging-us-central1-1077780374322-lwjldfuu/datacsv/part-0
0002-184cd7ba-1147-494d-ae8-e8da0bb34a25-c000.csv#1655742809770265...
Removing gs://dataproc-staging-us-central1-1077780374322-lwjldfuu/datacsv/part-0
0003-184cd7ba-1147-494d-ae8-e8da0bb34a25-c000.csv#1655742804954864...
/ [6 objects]
Operation completed over 6 objects.
```

```
22.18 GiB      gs://dataproc-staging-us-central1-1077780374322-lwjldfuu/datacsv/*
CPU times: user 172 ms, sys: 49.8 ms, total: 222 ms
Wall time: 6min 36s
```

## 2.1 Why do we care?

The compression of each data type matter when running different operations and computations, let's compare the 3

```
[21]: ## Create our dataframes
data_parquet = spark.read.parquet(f'{GCS_LOCATION}dataparquet')
data_csv = spark.read.csv(f'{GCS_LOCATION}datacsv')
```

```
data_avro = spark.read.format("avro").load(f'{GCS_LOCATION}/dataavro')
```

```
[22]: %%time
      data_parquet.count()
```

```
[Stage 11:=====> (12 + 2) / 14]
```

```
CPU times: user 8.26 ms, sys: 0 ns, total: 8.26 ms
```

```
Wall time: 2.24 s
```

```
[22]: 114420316
```

```
[23]: %%time
      data_csv.count()
```

```
[Stage 14:=====>(177 + 1) / 178]
```

```
CPU times: user 136 ms, sys: 42.3 ms, total: 178 ms
```

```
Wall time: 1min 2s
```

```
[23]: 114420316
```

```
[24]: %%time
      data_avro.count()
```

```
[Stage 17:=====> (37 + 1) / 38]
```

```
CPU times: user 67.6 ms, sys: 8.39 ms, total: 76 ms
```

```
Wall time: 2min
```

```
[24]: 114420316
```

- *What is the compression ratio for the parquet to csv file?* > We have  $1.7\text{G}/21\text{G} = 0.081$  or 8.1% of original size
- *Which serialization would query a column faster?* > Parquet has a columnar format therefore a column of data has faster access and only needs to grab a subset of data
- *Which types of columns do you think has the best compression for parquet?* > Columns with repeated content will have better compressions such as categorical columns will have very high compression ratios, especially if they're integers since parquet has enhanced compression for types with smaller storage requirements.
- *When should you use flat files vs other data formats?* > If you need human readable data or you have small data sets. Interoperability - for sharing with other teams. Don't send Bob in accounting a parquet file! Bob will try to open it in excel and he'll get an error, call IT, and IT will tell Bob to clear his cookies and restart his computer. Bob will not be impressed.

- *If we want to do analysis with lots of aggregations what serialization should we use?* > Parquet
- *Is there any downside to Parquet?* > Parquet is non-appendable (immutable) which means that if we have new data coming in we can't grow the dataset with parquet. Parquet datasets are typically used for batch analysis after the data has reached a final state, such as on a date roll-over.
- *If you had to partition data into days as new data comes in with aggregations happening at end of day how would you operationalize this?* > Data coming in for a day is streamed into an Avro file which handles appends seamlessly, then once the day has completed and a new partition for data is created a batch job can convert the avro file into a parquet file for the DS/Analyst team to query against.

### 3 Data Aggregation

Let's perform different aggregations using different methods and GroupBy. Don't worry! GroupBy from DataFrames is very different than RDDs.

```
[28]: %%time
      ## Let's start with sorting
      #data_parquet.sort("mean_temp").show()
      #data_parquet.sort("mean_temp").select("mean_temp").show()
      data_parquet.sort("mean_temp").select("mean_temp").filter(F.col("mean_temp").
      ↪isNotNull()).show()
```

[Stage 20:=====> (12 + 2) / 14]

```
+-----+
|      mean_temp|
+-----+
|      -118.5|
|-118.30000305175781|
| -117.5999984741211|
| -117.0999984741211|
| -115.9000015258789|
|-115.80000305175781|
|      -115.5|
|      -115.5|
|-115.19999694824219|
| -115.0999984741211|
|      -115.0|
| -114.9000015258789|
| -114.9000015258789|
| -114.5999984741211|
| -114.0999984741211|
|      -114.0|
| -113.9000015258789|
| -113.5999984741211|
| -113.5999984741211|
```



```
| -113.4000015258789|
+-----+
only showing top 20 rows
```

CPU times: user 19.8 ms, sys: 490 µs, total: 20.3 ms  
Wall time: 7.54 s

```
[29]: %%time
      ## Let's compare with avro
      data_avro.sort("mean_temp").select("mean_temp").filter(F.col("mean_temp").
      ↪isNotNull()).show()
```

[Stage 21:=====> (37 + 1) / 38]

```
+-----+
|      mean_temp|
+-----+
|      -118.5|
|-118.30000305175781|
| -117.5999984741211|
| -117.0999984741211|
| -115.9000015258789|
|-115.80000305175781|
|      -115.5|
|      -115.5|
|-115.19999694824219|
| -115.0999984741211|
|      -115.0|
| -114.9000015258789|
| -114.9000015258789|
| -114.5999984741211|
| -114.0999984741211|
|      -114.0|
| -113.9000015258789|
| -113.5999984741211|
| -113.5999984741211|
| -113.4000015258789|
+-----+
only showing top 20 rows
```

CPU times: user 63.2 ms, sys: 11.1 ms, total: 74.3 ms  
Wall time: 2min 1s

```
[30]: %%time
      data_parquet.select(F.mean("mean_wind_speed").alias("Avg mean Wind")).show()
```

[Stage 22:=====> (12 + 2) / 14]

```
+-----+
|    Avg mean Wind|
+-----+
|6.763403616672629|
+-----+
```

CPU times: user 7.71 ms, sys: 4.93 ms, total: 12.6 ms  
Wall time: 3.08 s

```
[31]: %%time
data_parquet.select(F.max("mean_wind_speed")).show()
```

[Stage 25:=====> (12 + 2) / 14]

```
+-----+
|max(mean_wind_speed)|
+-----+
|    96.9000015258789|
+-----+
```

CPU times: user 10.2 ms, sys: 1.07 ms, total: 11.3 ms  
Wall time: 2.5 s

```
[32]: %%time
data_parquet.select(F.min("mean_wind_speed")).show()
```

[Stage 28:=====> (11 + 3) / 14]

```
+-----+
|min(mean_wind_speed)|
+-----+
|                0.0|
+-----+
```

CPU times: user 9.27 ms, sys: 1.3 ms, total: 10.6 ms  
Wall time: 2.41 s

```
[33]: %%time
data_parquet.select(F.stddev("mean_wind_speed")).show()
```

[Stage 31:=====> (12 + 2) / 14]

```
+-----+
|stddev_samp(mean_wind_speed)|
```

```
+-----+
|                4.9270200129559|
+-----+
```


CPU times: user 10.5 ms, sys: 1.9 ms, total: 12.4 ms  
Wall time: 2.87 s

### 3.1 GroupBy

```
[34]: %%time
      ## Let's suppose we want the average temperature by year
      data_pandas = data_parquet.groupBy("year").agg(F.mean("mean_temp").alias("Mean_
      ↳Temp")).toPandas()
      data_pandas.sort_values("year").set_index("year").plot()
```

CPU times: user 2.38 s, sys: 163 ms, total: 2.54 s  
Wall time: 10.3 s

```
[34]: <AxesSubplot:xlabel='year'>
```



```
[35]: data_parquet.groupBy("year").agg(F.mean("mean_temp").alias("Mean Temp")).count()
```

[35]: 82


```
[36]: %%time
## Let's suppose we want more than one
data_pandas = data_parquet.groupBy("year").agg(F.mean("mean_temp").alias("Mean_Temp"),
                                              F.max("max_temperature").
                                              alias("Max Temp")).toPandas()
data_pandas.sort_values("year").set_index("year").plot()
```

[Stage 43:=====> (12 + 2) / 14]

CPU times: user 104 ms, sys: 34.1 ms, total: 138 ms

Wall time: 5.81 s

[36]: <AxesSubplot:xlabel='year'>



```
[37]: %%time
## Let's suppose we want more than one
data_pandas = data_parquet.groupBy("year").agg(F.mean("mean_temp").alias("Mean_Temp"),
                                              F.max("max_temperature").
                                              alias("Max Temp")).toPandas()
data_pandas.sort_values("year").set_index("year").plot()
```

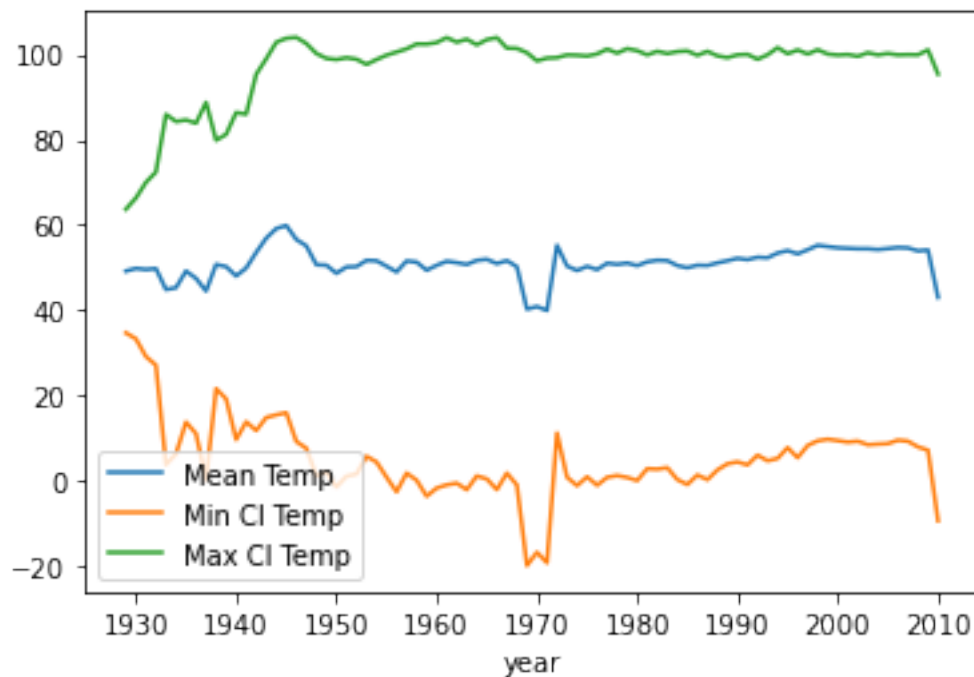
```

F.stddev("mean_temp").alias("SD_
↪Mean Temp")).toPandas()
data_pandas = data_pandas.sort_values("year").set_index("year")
data_pandas['Min CI Temp'] = data_pandas['Mean Temp'] - 2*data_pandas['SD Mean_
↪Temp']
data_pandas['Max CI Temp'] = data_pandas['Mean Temp'] + 2*data_pandas['SD Mean_
↪Temp']
del data_pandas['SD Mean Temp']
data_pandas.plot()

```

CPU times: user 130 ms, sys: 24.5 ms, total: 154 ms  
Wall time: 5.27 s

[37]: <AxesSubplot:xlabel='year'>



```

[38]: %%time
## Let's suppose we want more than one
data_parquet.groupBy(["year", 'month']).agg(F.mean("mean_temp").alias("Mean_
↪Temp"),
F.stddev("mean_temp").alias("SD_
↪Mean Temp")).sort(['year', 'month']).show()

```

[Stage 49:=====>

(12 + 2) / 14]

year	month	Mean Temp	SD Mean Temp
1929	8	60.552419416366085	3.1990462574283702
1929	9	61.32605038170053	4.033576425771579
1929	10	50.62124188117732	4.801964157921469
1929	11	47.099176249747615	6.103062771817114
1929	12	45.37819060909536	6.024023504130773
1930	1	44.1638498261501	4.924192609845936
1930	2	39.85389948800284	5.424807085849544
1930	3	43.270230277588496	6.20364872795214
1930	4	46.59947278336188	4.285578098496348
1930	5	51.23745814294719	5.8525031165383705
1930	6	57.698947384482935	4.381964851854148
1930	7	58.29831368123098	3.4151141395515636
1930	8	58.91847454006389	4.035375986521116
1930	9	57.15355978042948	5.05659012503392
1930	10	51.519968809464046	4.618773602848772
1930	11	46.13617694547391	6.240501759905245
1930	12	44.314175671899264	6.18088579805713
1931	1	40.092450479469676	7.572812511108157
1931	2	39.418145178466716	7.659416160405766
1931	3	40.05160875578184	8.428746256340379

only showing top 20 rows

CPU times: user 9.31 ms, sys: 7.83 ms, total: 17.1 ms  
Wall time: 6.19 s

### 3.2 User Defined Functions

```
[39]: ## Let's recall how we created the time column from before
data_parquet_time = data_parquet.withColumn("time",
                                             F.concat(F.col("year"),
                                             F.lit("-"), F.col("month"),
                                             F.lit("-"), F.col("day")) \
                                             .cast(types.TimestampType()))
```

```
[40]: %%time
data_parquet_time.select('time').show(5)
```

time
1929-10-20 00:00:00
1929-12-18 00:00:00

```
|1931-04-24 00:00:00|
|1931-06-23 00:00:00|
|1931-03-02 00:00:00|
+-----+
only showing top 5 rows
```

```
CPU times: user 3.42 ms, sys: 576 µs, total: 4 ms
Wall time: 483 ms
```

```
[41]: ## Can we do it differently? Yes! UDF. You can create UDF that will work row by
      ↳ row in your dataframe
def create_date_from_parts(year, month, day):
    return f'{year}-{month}-{day}'

create_date_udf = F.udf(create_date_from_parts, types.StringType())
data_parquet_time_udf = data_parquet.withColumn("time", create_date_udf('year',
↳ 'month', 'day')).cast(types.TimestampType())
```

```
[42]: %%time
data_parquet_time_udf.select('time').show(5)
```

```
[Stage 53:> (0 + 1) / 1]
```

```
+-----+
|          time|
+-----+
|1929-10-20 00:00:00|
|1929-12-18 00:00:00|
|1931-04-24 00:00:00|
|1931-06-23 00:00:00|
|1931-03-02 00:00:00|
+-----+
only showing top 5 rows
```

```
CPU times: user 5.98 ms, sys: 0 ns, total: 5.98 ms
Wall time: 1.49 s
```

UDFs are typically much slower than built-in Spark functionality. The reason for this is because they have to serialize and deserialize the data for every row that the function is applied to. There have been recent improvements to UDF for some analytical results with Pandas UDFs that return scalars or groupby maps. Some more information about why UDFs are inefficient can be found here <https://blog.cloudera.com/blog/2017/02/working-with-udfs-in-apache-spark/>

```
[43]: %%time
      ## Let's look at other examples
from pyspark.sql.functions import udf
@udf("double")
```

```
def squared_udf(s):
    return s * s

data_udf = data_parquet.withColumn("square_temp", squared_udf(F.
    ↪col("mean_temp")))
data_udf.select("square_temp").show()
```

```
+-----+
|      square_temp|
+-----+
|2787.8399194335943|
|          2256.25|
|2520.0400765991217|
| 4238.009801330569|
|1831.8399346923834|
|          4489.0|
| 4678.560208740237|
| 4108.809804382327|
|1689.2098745727562|
| 605.1600187683107|
|          5041.0|
|          5041.0|
| 3047.040084228516|
| 3058.089915618897|
|          121.0|
| 3931.290095672608|
|          4970.25|
|          25.0|
|106.09000392913822|
|          4225.0|
+-----+
```

only showing top 20 rows

CPU times: user 9.54 ms, sys: 3.68 ms, total: 13.2 ms

Wall time: 741 ms

```
[44]: %%time
data_no_udf = data_parquet.withColumn("square_temp", F.col("mean_temp")**2)
data_udf.select("square_temp").show()
```

```
+-----+
|      square_temp|
+-----+
|2787.8399194335943|
|          2256.25|
|2520.0400765991217|
```



```
| 4238.009801330569|
|1831.8399346923834|
|          4489.0|
| 4678.560208740237|
| 4108.809804382327|
|1689.2098745727562|
| 605.1600187683107|
|          5041.0|
|          5041.0|
| 3047.040084228516|
| 3058.089915618897|
|          121.0|
| 3931.290095672608|
|          4970.25|
|          25.0|
|106.09000392913822|
|          4225.0|
```

```
+-----+
```

only showing top 20 rows

CPU times: user 5.16 ms, sys: 0 ns, total: 5.16 ms

Wall time: 522 ms

```
[45]: %%time
## You can also use UDF with select
data_parquet.select("mean_temp", squared_udf("mean_temp").
    ↪alias("squared_temp")).show()
```

```
+-----+
|          mean_temp|          squared_temp|
+-----+
| 52.79999923706055|2787.8399194335943|
|          47.5|          2256.25|
| 50.20000076293945|2520.0400765991217|
| 65.0999984741211| 4238.009801330569|
| 42.79999923706055|1831.8399346923834|
|          67.0|          4489.0|
| 68.4000015258789| 4678.560208740237|
| 64.0999984741211| 4108.809804382327|
| 41.099998474121094|1689.2098745727562|
| 24.600000381469727| 605.1600187683107|
|          71.0|          5041.0|
|          71.0|          5041.0|
| 55.20000076293945| 3047.040084228516|
| 55.29999923706055| 3058.089915618897|
|          -11.0|          121.0|
| 62.70000076293945| 3931.290095672608|
|          70.5|          4970.25|
```

	-5.0	25.0
-10.300000190734863	106.09000392913822	
	65.0	4225.0

only showing top 20 rows

CPU times: user 0 ns, sys: 5.41 ms, total: 5.41 ms  
Wall time: 459 ms

## 4 Spark SQL

Finally, let's work with Spark SQL. Spark allows us to combine the power of SQL with Spark and the Dataframes API

```
[46]: %%time
      ## Let's run an example
      # First we need to create a temporary table that we can query
      data_parquet.registerTempTable('data')
      spark.sql(
        """
        select mean_temp
        from data
        """).show()
```

mean_temp
52.79999923706055
47.5
50.20000076293945
65.0999984741211
42.79999923706055
67.0
68.4000015258789
64.0999984741211
41.099998474121094
24.600000381469727
71.0
71.0
55.20000076293945
55.29999923706055
-11.0
62.70000076293945
70.5
-5.0
-10.300000190734863
65.0

only showing top 20 rows

CPU times: user 0 ns, sys: 4 ms, total: 4 ms

Wall time: 554 ms

```
[47]: %%time
data_parquet.select("mean_temp").show()
```

```
+-----+
|          mean_temp|
+-----+
| 52.79999923706055|
|              47.5|
| 50.20000076293945|
| 65.0999984741211|
| 42.79999923706055|
|              67.0|
| 68.4000015258789|
| 64.0999984741211|
| 41.099998474121094|
| 24.600000381469727|
|              71.0|
|              71.0|
| 55.20000076293945|
| 55.29999923706055|
|              -11.0|
| 62.70000076293945|
|              70.5|
|              -5.0|
|-10.300000190734863|
|              65.0|
+-----+
```

only showing top 20 rows

CPU times: user 0 ns, sys: 4.84 ms, total: 4.84 ms

Wall time: 369 ms

```
[48]: ## Let's run multiple querys similar to the ones we ran before
spark.sql(
    """
    select mean_temp, power(mean_temp, 2) as squared_temp
    from data
    """).show()
```

```
+-----+-----+
|          mean_temp|    squared_temp|
+-----+-----+
| 52.79999923706055|2787.8399194335943|
```

	47.5	2256.25
50.20000076293945	2520.0400765991217	
65.0999984741211	4238.009801330569	
42.79999923706055	1831.8399346923834	
	67.0	4489.0
68.4000015258789	4678.560208740237	
64.0999984741211	4108.809804382327	
41.099998474121094	1689.2098745727562	
24.600000381469727	605.1600187683107	
	71.0	5041.0
	71.0	5041.0
55.20000076293945	3047.040084228516	
55.29999923706055	3058.089915618897	
	-11.0	121.0
62.70000076293945	3931.290095672608	
	70.5	4970.25
	-5.0	25.0
-10.300000190734863	106.09000392913822	
	65.0	4225.0

+-----+

only showing top 20 rows

```
[49]: %%time
## Let's run multiple queries similar to the ones we ran before
spark.sql(
    """
    select
        year,
        month,
        avg(mean_temp) as mean,
        std(mean_temp) as st_dev
    from
        data
    group by
        year,
        month
    order by
        year,
        month
    """).show()
```

[Stage 60:=====>

(12 + 2) / 14]

year	month	mean	st_dev
1929	8	60.552419416366085	3.1990462574283702

1929	9	61.32605038170053	4.033576425771579
1929	10	50.62124188117732	4.801964157921469
1929	11	47.099176249747615	6.103062771817114
1929	12	45.37819060909536	6.024023504130773
1930	1	44.1638498261501	4.924192609845936
1930	2	39.85389948800284	5.424807085849544
1930	3	43.270230277588496	6.20364872795214
1930	4	46.59947278336188	4.285578098496348
1930	5	51.23745814294719	5.8525031165383705
1930	6	57.698947384482935	4.381964851854148
1930	7	58.29831368123098	3.4151141395515636
1930	8	58.91847454006389	4.035375986521116
1930	9	57.15355978042948	5.05659012503392
1930	10	51.519968809464046	4.618773602848772
1930	11	46.13617694547391	6.240501759905245
1930	12	44.314175671899264	6.18088579805713
1931	1	40.092450479469676	7.572812511108157
1931	2	39.418145178466716	7.659416160405766
1931	3	40.05160875578184	8.428746256340379

+---+---+-----+-----+

only showing top 20 rows

CPU times: user 11.2 ms, sys: 3.23 ms, total: 14.4 ms

Wall time: 6.03 s

```
[50]: ## We can save the Spark SQL query as a dataframe
df_sql = spark.sql(
    """
    select
        year,
        month,
        avg(mean_temp) as mean,
        std(mean_temp) as st_dev
    from
        data
    group by
        year,
        month
    order by
        year,
        month
    """)

df_sql_pd = df_sql.toPandas()
df_sql_pd = df_sql_pd.set_index(["year", 'month'])
df_sql_pd
```

```
[50]:
```

		mean	st_dev
year	month		
1929	8	60.552419	3.199046
	9	61.326050	4.033576
	10	50.621242	4.801964
	11	47.099176	6.103063
	12	45.378191	6.024024
...		...	...
2009	12	39.834150	26.892829
2010	1	37.385471	28.259400
	2	39.672241	27.192215
	3	46.851747	23.391840
	4	53.084282	20.170773

[969 rows x 2 columns]

```
[52]: %%time
## Let's also join dataframes
stations = spark.read.format('bigquery') \
    .option('table', 'bigquery-public-data:noaa_gsod.stations') \
    .load()

stations_us = stations.filter(F.col('Country')== 'US')

## One of the dataframes is quite small, so let's broadcast it!
# join_data = data_parquet.join(F.broadcast(stations_us), stations_us.
    ↳ usaf==data_parquet.station_number, 'inner')
join_data = data_parquet.join(stations_us, stations_us.usaf==data_parquet.
    ↳ station_number, 'inner')
join_data.count()
```

[Stage 76:=====> (12 + 2) / 14]

CPU times: user 17.1 ms, sys: 4.96 ms, total: 22 ms

Wall time: 19.2 s

[52]: 4584888375

```
[53]: join_data.printSchema()
```

```
root
 |-- station_number: long (nullable = true)
 |-- wban_number: long (nullable = true)
 |-- year: long (nullable = true)
 |-- month: long (nullable = true)
```

```

|-- day: long (nullable = true)
|-- mean_temp: double (nullable = true)
|-- num_mean_temp_samples: long (nullable = true)
|-- mean_dew_point: double (nullable = true)
|-- num_mean_dew_point_samples: long (nullable = true)
|-- mean_sealevel_pressure: double (nullable = true)
|-- num_mean_sealevel_pressure_samples: long (nullable = true)
|-- mean_station_pressure: double (nullable = true)
|-- num_mean_station_pressure_samples: long (nullable = true)
|-- mean_visibility: double (nullable = true)
|-- num_mean_visibility_samples: long (nullable = true)
|-- mean_wind_speed: double (nullable = true)
|-- num_mean_wind_speed_samples: long (nullable = true)
|-- max_sustained_wind_speed: double (nullable = true)
|-- max_gust_wind_speed: double (nullable = true)
|-- max_temperature: double (nullable = true)
|-- max_temperature_explicit: boolean (nullable = true)
|-- min_temperature: double (nullable = true)
|-- min_temperature_explicit: boolean (nullable = true)
|-- total_precipitation: double (nullable = true)
|-- snow_depth: double (nullable = true)
|-- fog: boolean (nullable = true)
|-- rain: boolean (nullable = true)
|-- snow: boolean (nullable = true)
|-- hail: boolean (nullable = true)
|-- thunder: boolean (nullable = true)
|-- tornado: boolean (nullable = true)
|-- usaf: string (nullable = true)
|-- wban: string (nullable = true)
|-- name: string (nullable = true)
|-- country: string (nullable = true)
|-- state: string (nullable = true)
|-- call: string (nullable = true)
|-- lat: double (nullable = true)
|-- lon: double (nullable = true)
|-- elev: string (nullable = true)
|-- begin: string (nullable = true)
|-- end: string (nullable = true)

```

```

[54]: %%time
## let's now use Spark SQL to query this data
join_data.registerTempTable('data')
spark.sql(
    """
select
    state,

```

```

    avg(mean_temp) as mean,
    avg(lat),
    avg(lon)
from
    data
where
    state in ('CA', 'TX', 'NY')
group by
    state
order by
    state
""").show()

```

[Stage 80:=====> (13 + 1) / 14]

state	mean	avg(lat)	avg(lon)
CA	53.747206531956735	35.89247422036925	-119.49590371314966
NY	53.66972249742855	42.22919509172085	-76.01350616411186
TX	53.81268076536781	30.791588005852184	-98.54278344775695

CPU times: user 27.7 ms, sys: 0 ns, total: 27.7 ms

Wall time: 43 s

[ ]: