

Streaming Systems

Streaming Systems: Objectives

1. Theory

- Define:
 - Data latency
 - Windowing
 - Watermarking
 - Triggers
 - Bounded data
- Understand exactly-once vs. at-least-once semantics
- Understand how streaming systems handle late data
- Understand which windowing strategy should be used based on use case

Streaming Systems: Objectives

2. Practice

- Lambda/kappa architecture
- Articulate when a use case requires streaming architectures

Streaming Systems: Table of Contents

1. Data latency
2. Streaming vs. batch
3. Why streaming?
4. Architectures
 - a. Lambda architecture
 - b. Kappa architecture
5. Exactly-once vs. at-least-once
6. Windowing
7. Additional considerations
8. Use cases
 - a. Ad monetization
 - b. Change data capture

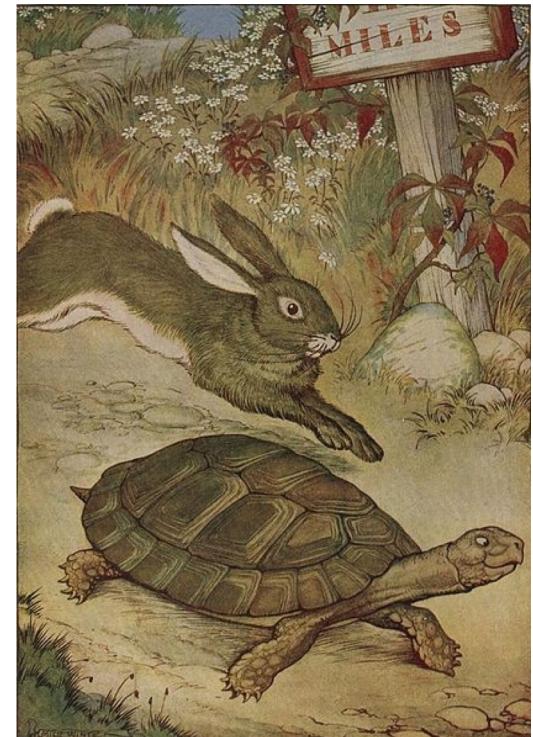
Streaming Systems

The End

Data Latency

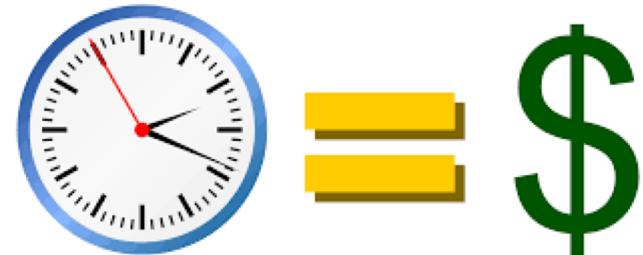
Real-Time, Near-Time, Some-Time

- Real-time
 - We need the answer as part of the process to process an event.
- Near-time
 - We need the answer shortly after an event (less than 10 seconds or less than five minutes, based on use case).
- Some-time
 - We need the answer in the future (hourly, daily, monthly, yearly).



Real-Time, Near-Time, Some-Time Concerns

- Real-time systems are much more expensive than some-time systems.
- Determine the criteria for when to use a real-time system to justify the expense.
- Real-time systems also carry additional complexity, as we will discuss.



Why Do We Care About Streaming?

- Having the answers earlier can provide an advantage in many ways.
 - Efficiency
 - Strategically
 - Responsiveness
- When we have a need for streaming systems, it becomes painful to deal with the batch systems.
- Did that update affect my users?
 - If so, roll it back
- Is my brand under attack on social media?
 - Get PR involved earlier
- Is there a good deal on the stock market?
 - Buy/sell

Data Latency

The End

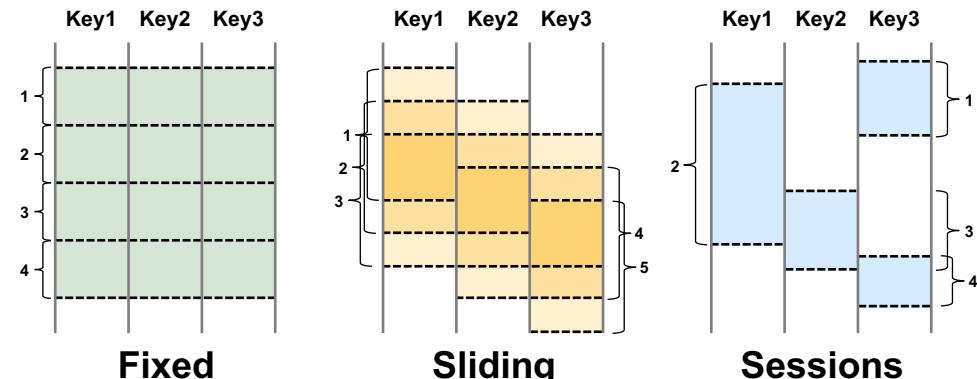
Streaming vs. Batch

Bounded vs. Unbounded Data

- Bounded data have a fixed volume to process
- Theoretically, all batch workload
 - Data might have to be split to be able to be processed
 - Directly in line with MapReduce
- Unbounded data do not have a fixed volume
- Practically all data are bounded
 - Can't have infinite data
 - Focus on the volume of data that the system needs to handle

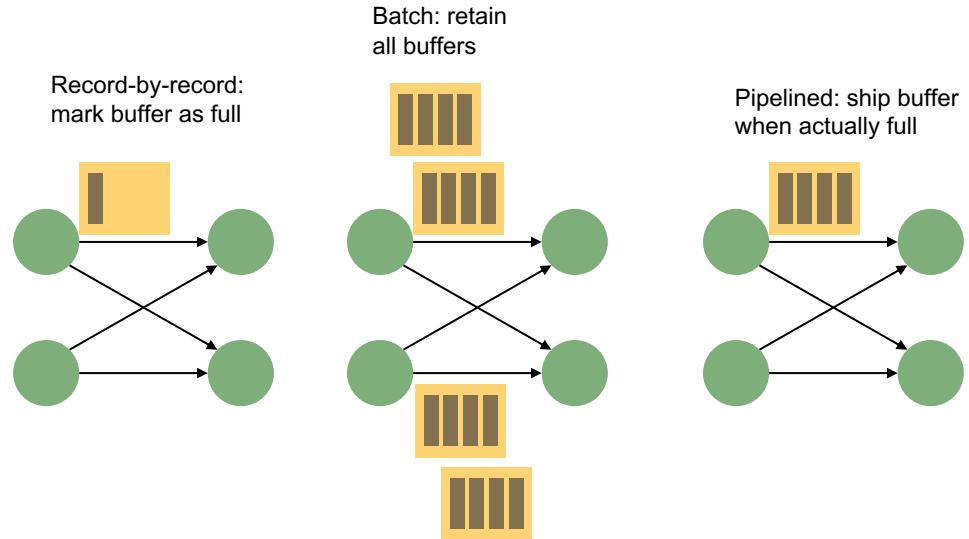
Windows

- Chop up the data based on a time interval such as event or processing time
 - More on this later
- Streaming systems are designed to deal with complexities that arise in windowing when it comes to completeness of data
- A few ways to window data
 - Fixed
 - Sliding
 - Sessions



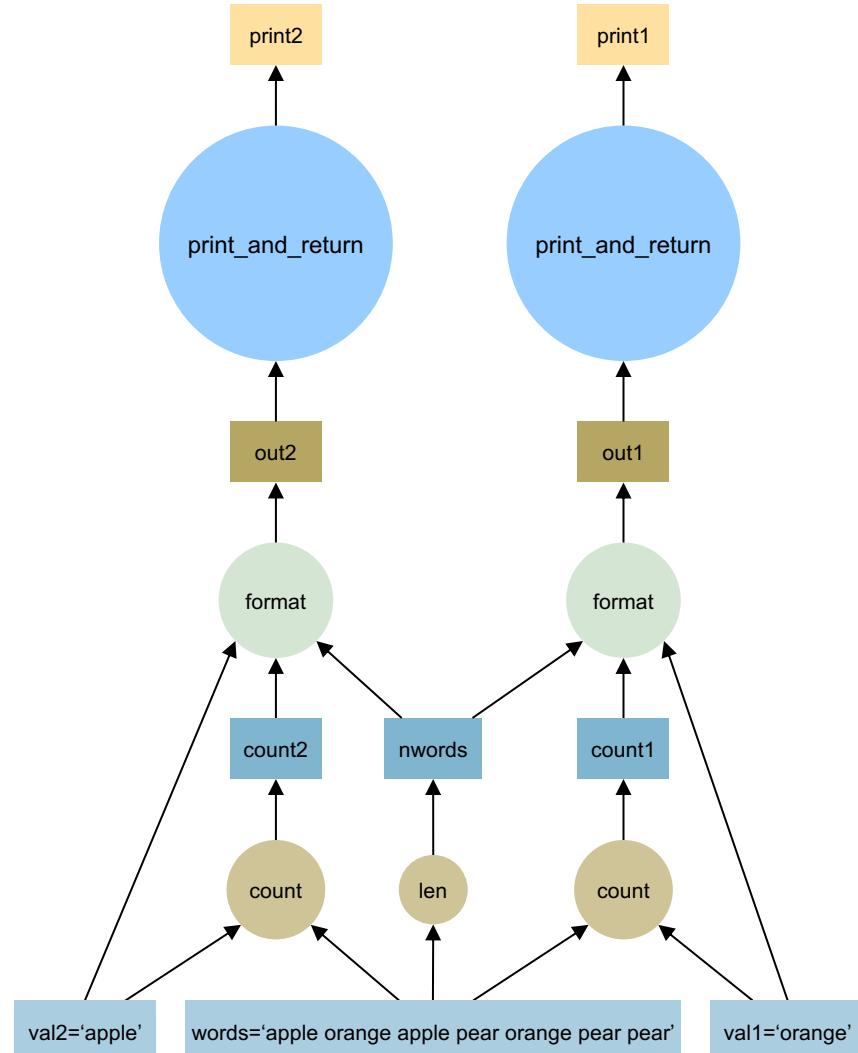
Runtimes

- When do we ship data for processing?
- There are several options
 - Record-by-record
 - Pipelined
 - Batch
- Each has their own advantages, and different systems implement management of these runtimes differently



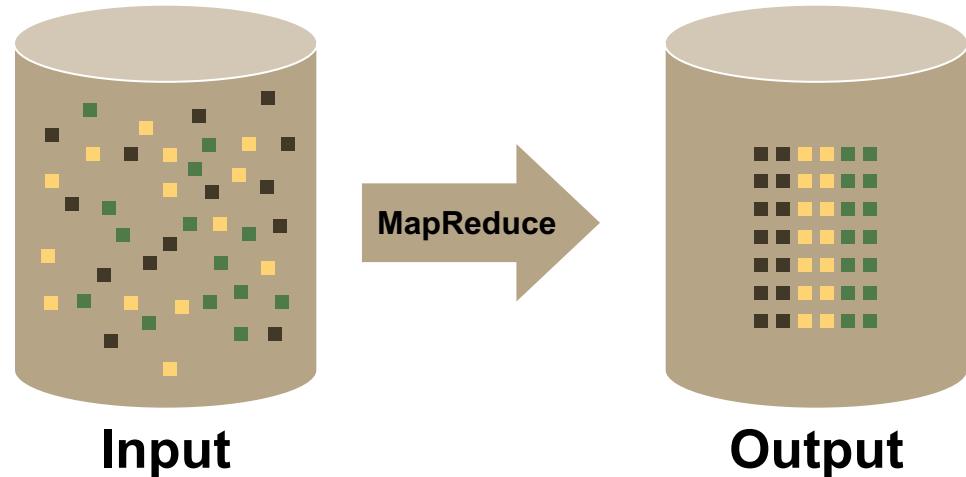
DAGs

- Directed acyclic graphs (DAGs) are sequences of events that need to follow each other that allow for splitting of independent work as we've seen before
- Batch DAGs run each step that's not blocked as it finishes processing
- Streaming DAGs takes this one step further and pushes a given data buffer through the DAG
- Compliments the runtimes we just discussed
- Every batch process starts by streaming in a file



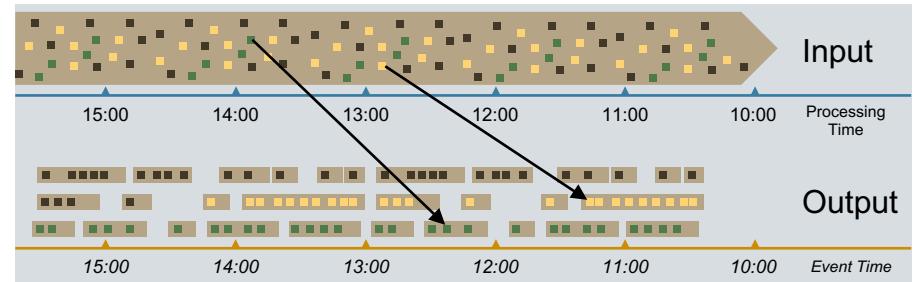
What Is Batch Processing?

- Bounded data
- Global window
 - Can be broken up into fixed windows
- “Batch” runtime
 - Retain all buffers of data
- Blocking DAGs



What Is Stream Processing?

- Unbounded data
- Nonglobal windows
 - Fixed
 - Sliding
 - Session
- “Record-by-record” runtime
 - Can also do micro-batching with “pipelined” runtimes
- Streaming DAGs
 - Batch does this in disguise while reading inputs



Streaming vs. Batch

The End

Architectures

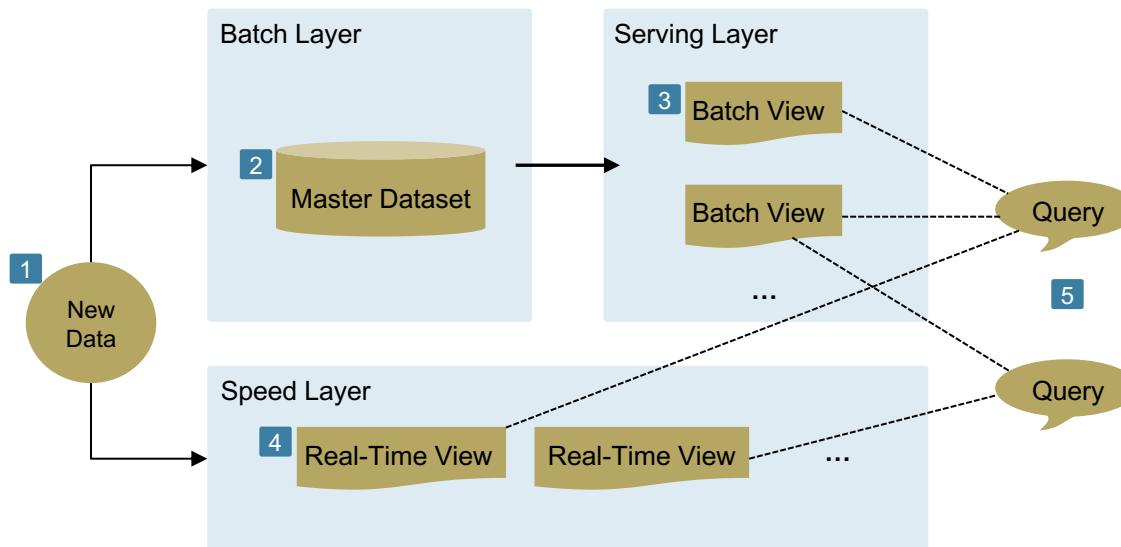
Why Architecture?

- To make a system work, we need to think about how we should put it in place
- Streaming systems typically follow one of two approaches:
 1. Lambda architecture
 2. Kappa architecture
- These approaches have different trade-offs and have different advantages associated with them
- A good system is built on a solid foundation



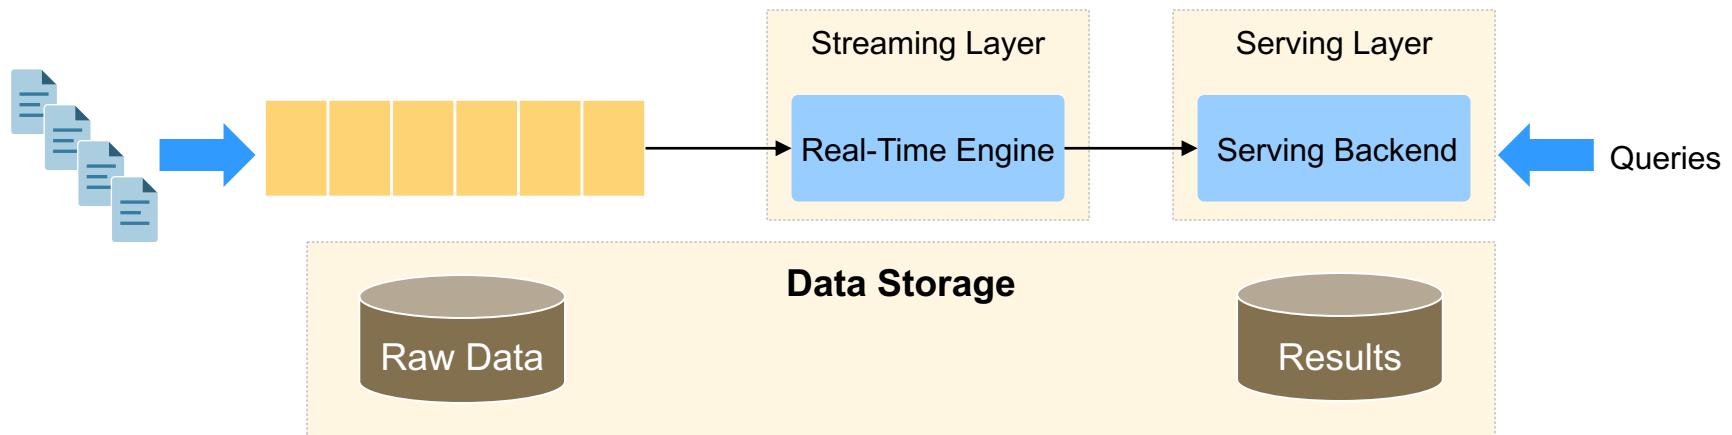
Lambda Architecture

- Manage two systems—one batch and one streaming.
 - A lot of code to manage
- Streaming systems historically did not handle missing or late data very effectively.



Kappa Architecture

- The world is a stream of data
- Offload data to batch for when needed
- Serve out of the streaming layer
- Much simpler to manage
- How do you deal with the gotcha?
 - Reconciliation process for late/missing data



Exactly-Once vs. At-Least-Once Processing

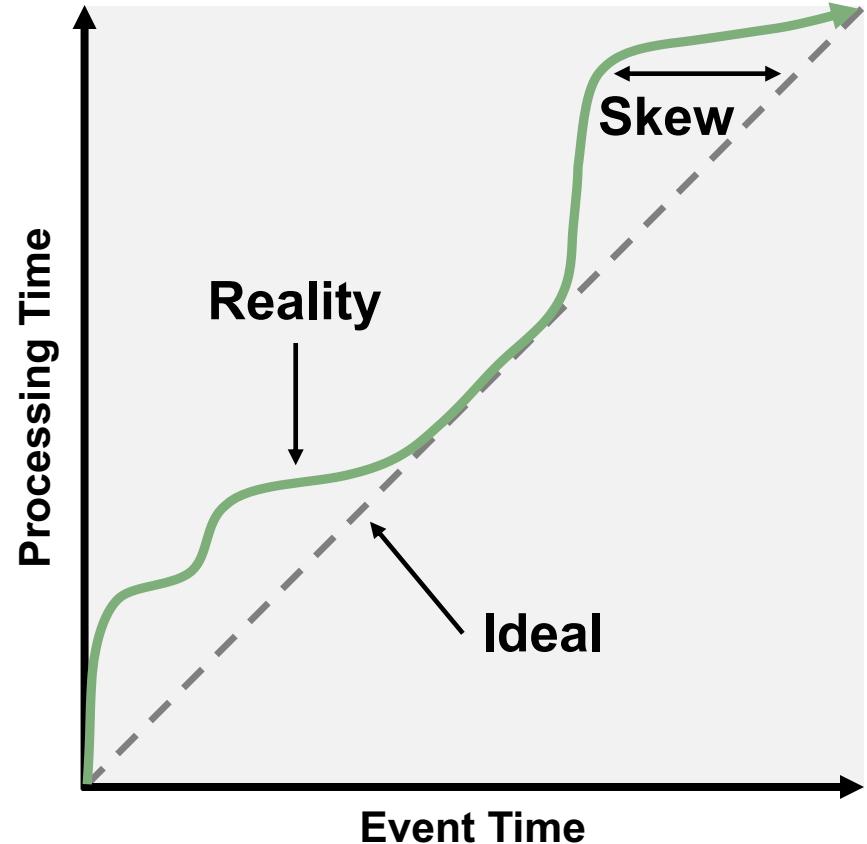
Understanding Streaming Semantics



At most once	At least once	Exactly once
Message pulled once	Message pulled one or more times; processed each time	Message pulled one or more times; processed once
May or may not be received	Receipt guaranteed	Receipt guaranteed
No duplicates	Likely duplicates	No duplicates
Possible missing data	No missing data	No missing data

Real-World Performance Considerations

- No system has perfect communication.
- Latency will always be a factor.
- What if a user walks into an elevator?
- How do we handle missing data?



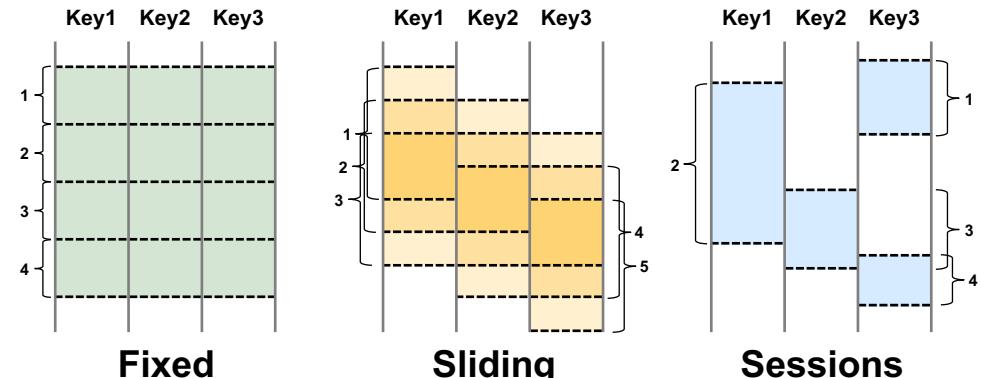
Architectures

The End

Considerations

Windowing

- Fixed
 - GroupBy
- Sliding
 - Look back five minutes from end of each window
 - Tracking a metric over time such as moving average of financial transactions
- Sessions
 - Dynamic window example
 - Typically used for user behavior
 - Trigger to end session when no messages for some fixed period



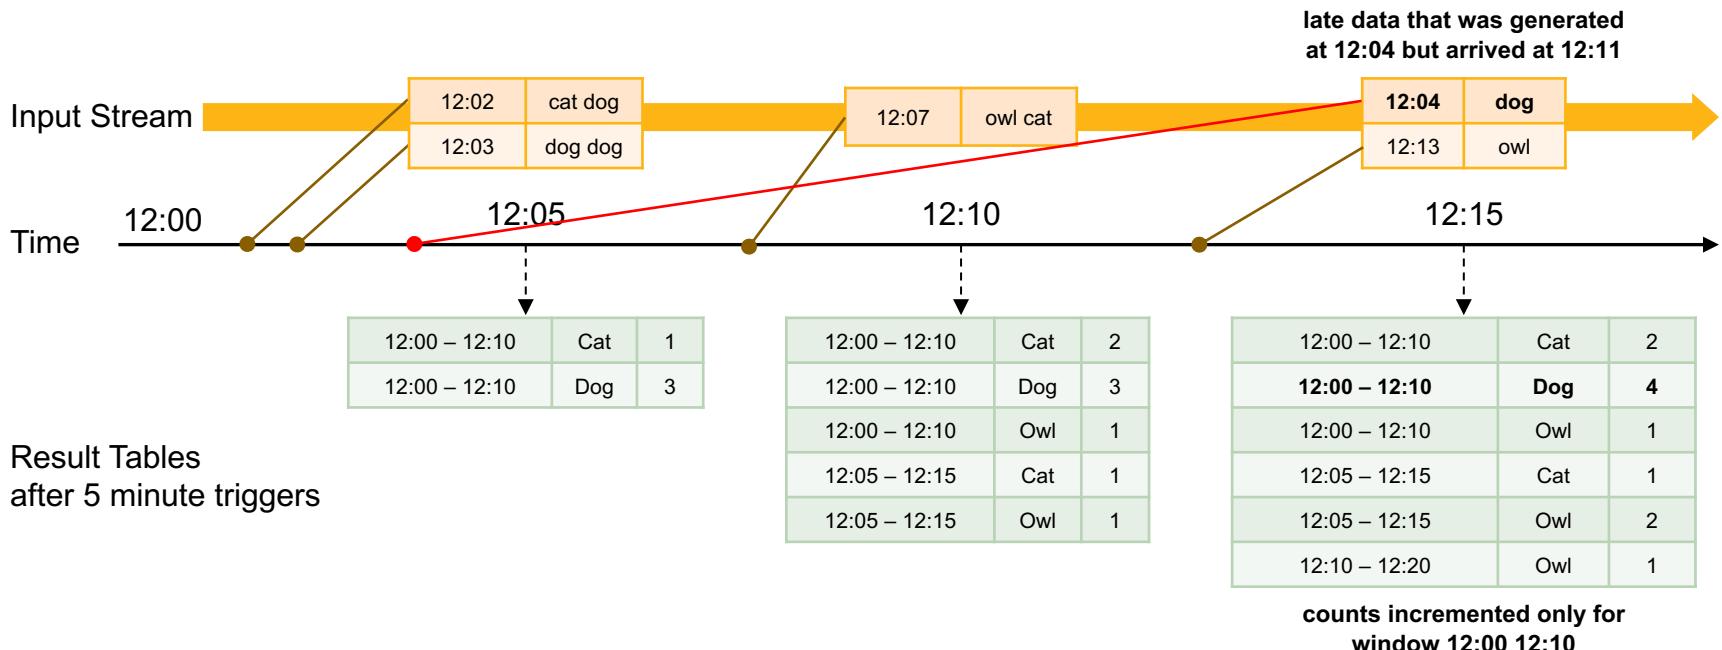
Triggers

- When to do the action
- Allows for us to view the output of a window again
- Typical built-in triggers include:
 - Count above N
 - Event time
 - Processing time
 - Run one buffer
- Different systems have different semantics, and the triggers will be based on those



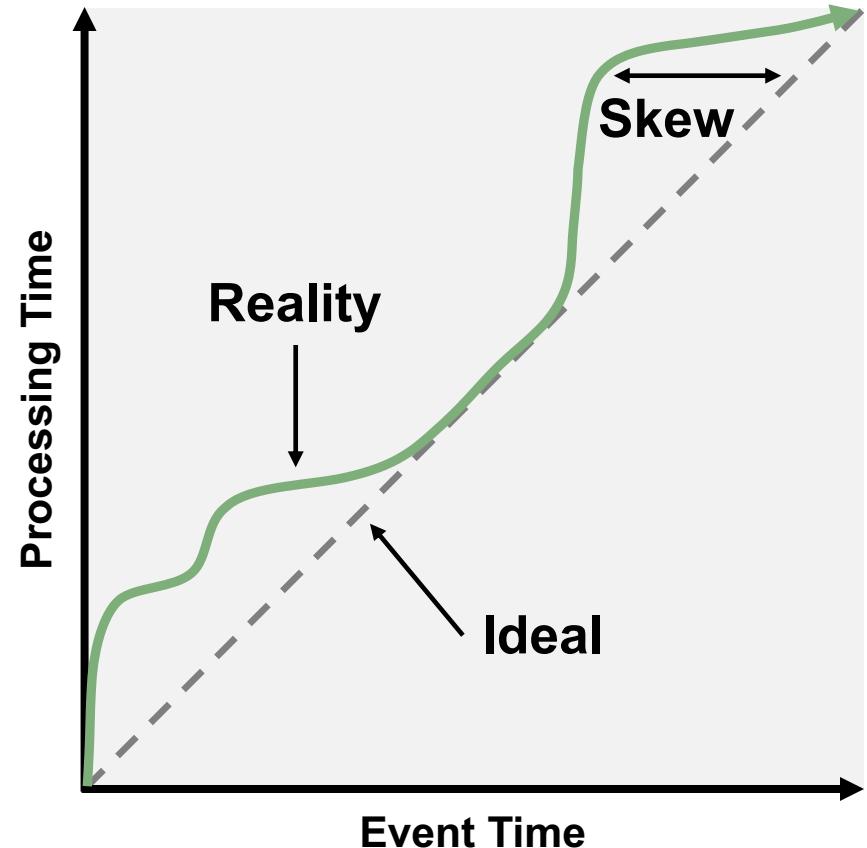
Late Data

- Sometimes data can show up late
- The system needs to be able to handle this
- Example of a fixed window aggregation and five-minute triggers with late data:



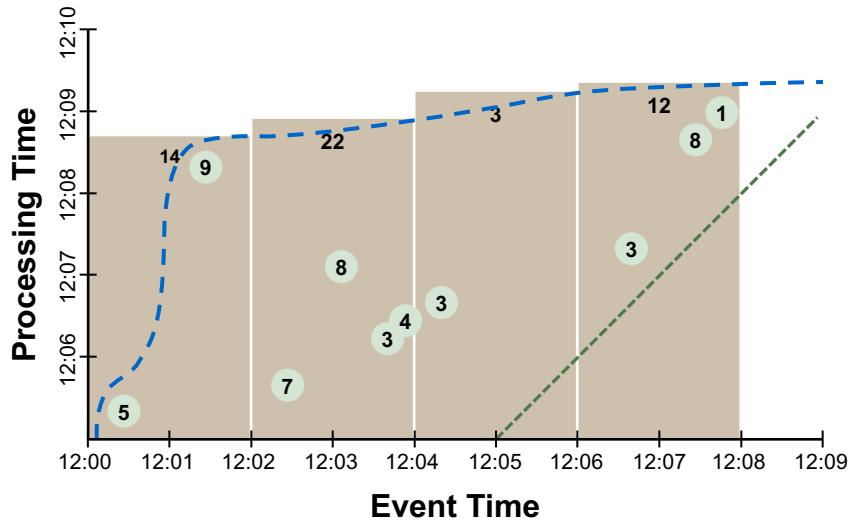
Watermarking

- We need to be able to keep track of where we are in the processing.
- The world is cruel, and we never have a perfect system.
- Watermarking takes a point in processing time and maps it to an event time.
 - It would be the red line in the figure.
- A perfect watermark would be the diagonal where processing time is equal to event time.
- Data that come in late would likely fall above the red line.

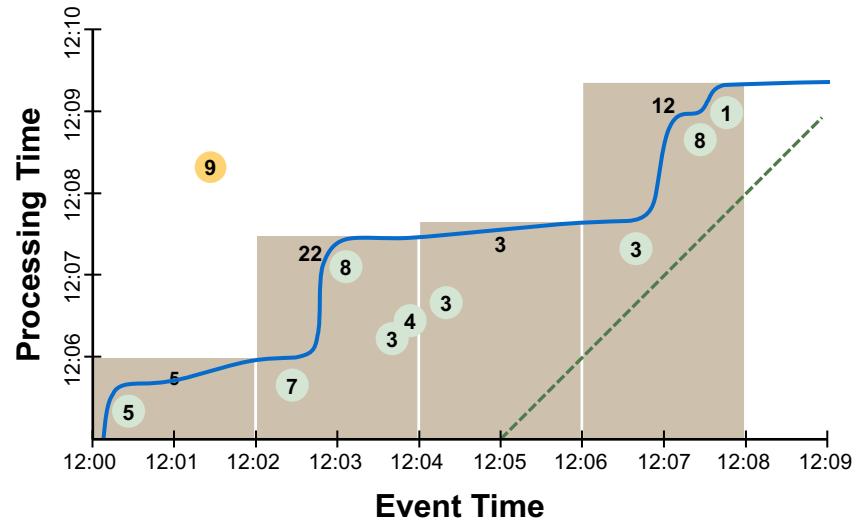


Watermarking

Perfect Watermark



Heuristic Watermark



Perfect watermark:

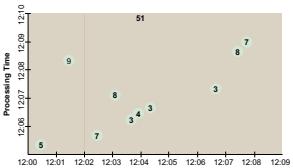
Ideal watermark:

Heuristic watermark:

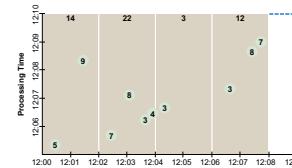
Ideal watermark:

Going Forward

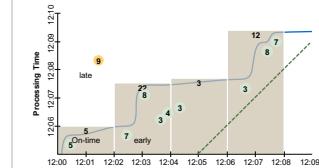
- This class is meant for an introduction to streaming.
 - If you want to get deeper into these ideas, review the following as a first step:
 - Event time vs. processing time for watermarking and windowing
 - Perfect vs. heuristic watermarking



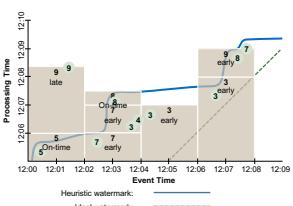
Classic batch



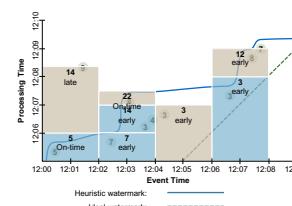
Fixed windows batch



Fixed windows streaming watermark



Early/late discarding

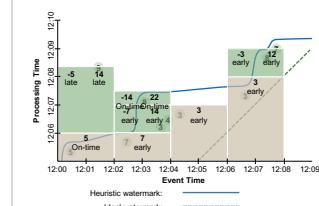


Early/late discarding

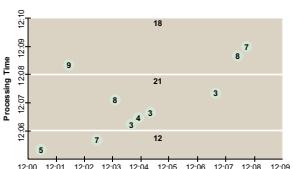
Listing 7 / Figure 9

Early/late accumulation

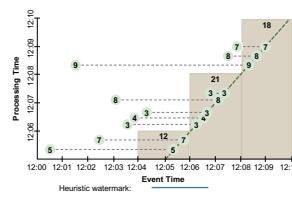
Listing 4 and 5 / Figure 10



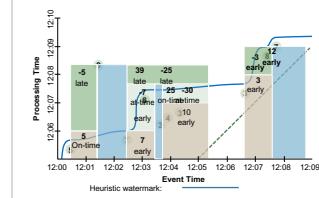
Early/late retracting



Processing – time (triggers)



Listing 10 / Figure 15



Sessions

Considerations

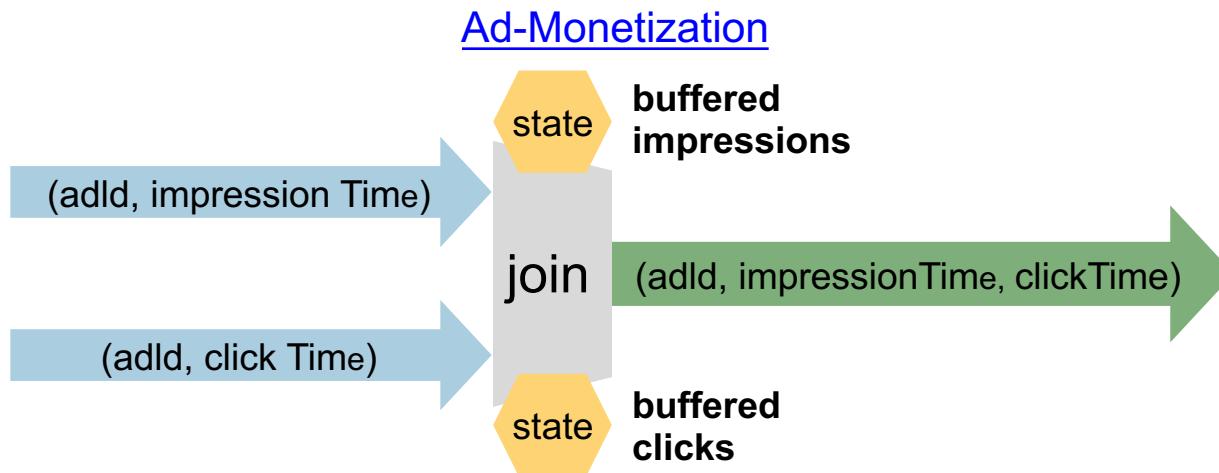
The End

Use Cases

Ad Monetization

- Advertisers want information now to ensure that their campaigns are attracting customers and generating revenue
- Two streams of data:
 1. Ads
 2. Clicks
- Need to join the streams, but there are several considerations

The Case for Stream-Stream Joins:

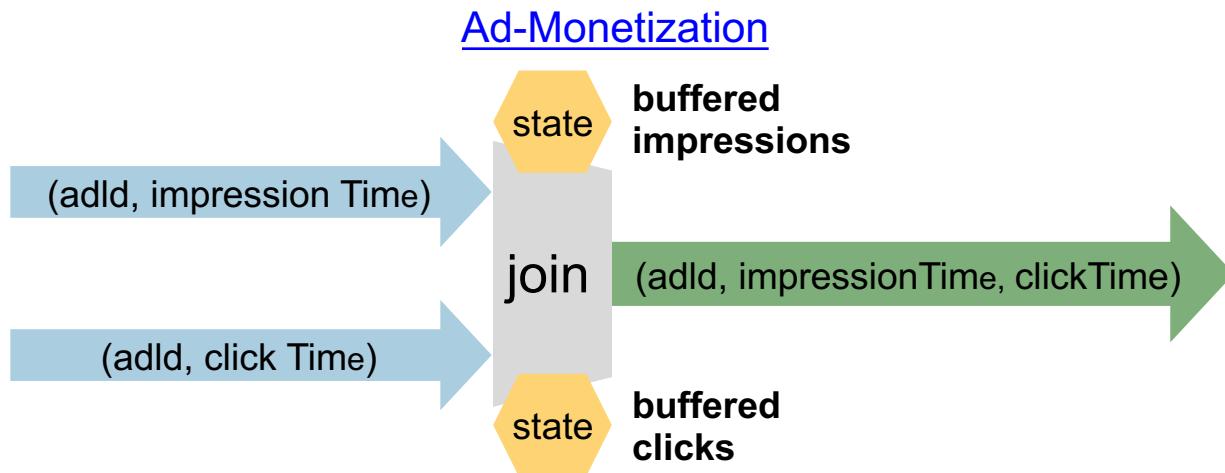


Stream-stream join use case: Ad-Monetization (joining ad clicks to impressions)

Ad Monetization

- What if there never was a click?
- How long should I wait for a click?
- What are the infrastructure timings?

The Case for Stream-Stream Joins:



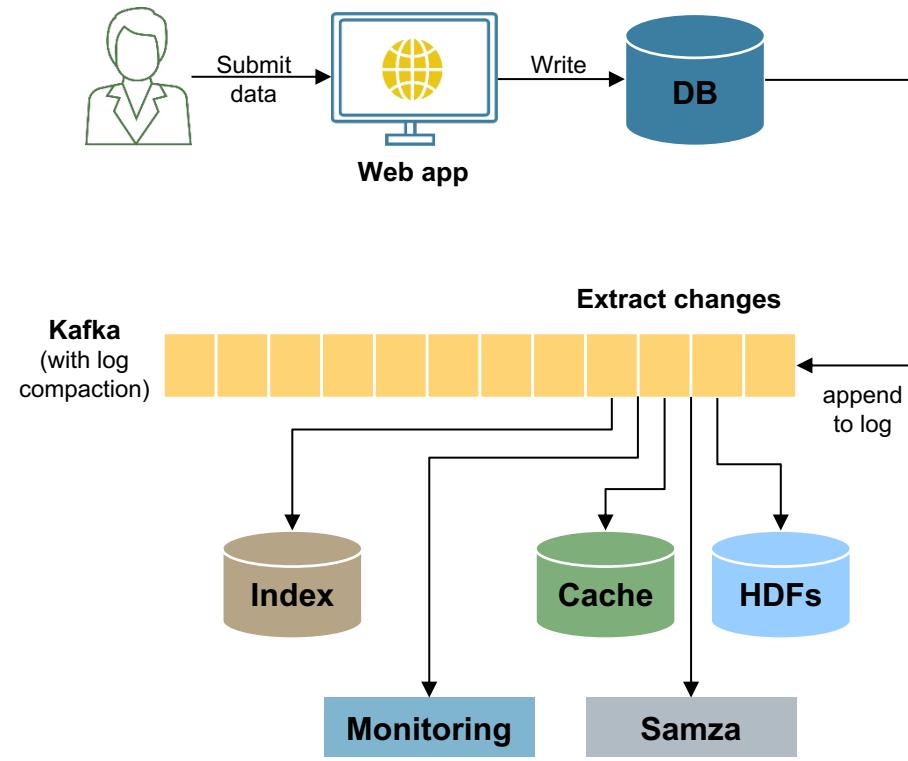
Stream-stream join use case: Ad-Monetization (joining ad clicks to impressions)

Change Data Capture

- You have a database on premises that is very sensitive to nonapplication workloads running against it
- To enable analytics against the application data, you need to replicate these data somewhere else
- Creating a stream of transactions committed to the database allows you to replicate the database somewhere else, and it is called change data capture (CDC)
- Single stream and one of the oldest use cases for streaming data

How to Integrate Your Databases

Using Change Capture



Use Cases

The End

Summary

Conclusion

- We won't be experts in a day.
- Streaming can be complicated, so limit its use to where it makes sense.
 - Additionally, streaming systems are more expensive than batch systems.
- Understand the basics of watermarking and how it relates to event time vs. processing time.
- Understand the semantics to consider for exactly-once vs. at-least-once processing.
- Speak to example use cases that involve streaming systems.



Streaming Systems: Outcomes

1. Theory

- Define
 - Data latency
 - Windowing
 - Watermarking
 - Triggers
 - Bounded data
- Understand exactly-once vs. at-least-once semantics
- Understand how streaming systems handle late data
- Understand which windowing strategy should be used based on use case

Streaming Systems: Outcomes

2. Practice

- Lambda/kappa architecture
- Articulate when a use case requires streaming architectures

Summary

The End